



Governing Artificial Intelligence: Law, Policy, and Institutions

Law 4052 // Computer Science 283 // Communication 252A // Political Science 145B/445B
Fall Quarter 2026




Rob Reich
Anka Reuel
Nate Persily


AI Governance: A Multi-Disciplinary Exploration


- Integrate philosophy/geopolitics with law/policy and technical approaches
 - Integrated assignments
- Why?
 - Global problems do not come neatly sorted into clean disciplinary boxes
- Historical illustration: the nuclear age
 - Physicists invent the nuclear age; game theorists manage its geopolitics


What AI Governance is NOT



AI Law and Regulation

A large, bold red 'X' is superimposed over the text 'AI Law and Regulation', effectively crossing it out. The 'X' is composed of two thick, slightly irregular red lines that intersect in the center. The text is in a black, sans-serif font and is partially obscured by the 'X'.





AI Governance


1. Law, policy, and regulation that emanates distinctively from government(s)

+

2. Professional norms, standards, tooling, and responsibilities that develop from within the respective profession

About me: a philosopher's path

Trends in Computer Science Degrees Over the 2010s



The Atlantic

TECHNOLOGY

Stanford's Top Major Is Now Computer Science

THE
NEW YORKER

GET RICH U.

There are no walls between Stanford and Silicon Valley. Should there be?

GOOD INTENTIONS

STANFORD TAKES ON THE TECHLASH

With more and more students becoming dorm-room C.E.O.s, three professors cooked up an ethics class for the coding set.



By Andrew Marantz

October 18, 2021



In the fall of 2015, Rob Reich, a philosopher and a political scientist at Stanford, was chatting with a freshman during office hours. “I asked him what he planned to study,”



"BRILLIANTLY ARGUED...ESSENTIAL READING FOR ANYONE WORRIED
ABOUT MONEY IN POLITICS."—LARISSA MACFARQUHAR

JUST GIVING WHY PHILANTHROPY IS FAILING DEMOCRACY AND HOW IT CAN DO BETTER ROB REICH

"SPELLS OUT
WHAT NEEDS
TO BE FIXED."*
—Wall Street Journal

"READ THIS."**
—Reed Hastings,
CEO of Netflix

"PROFOUNDLY
IMPORTANT."***
—Dr. Fei-Fei Li

"BRILLIANTLY
CRAFTED."

—Evan Spiegel,
CEO of SnapChat

"A TRIUMPH."


—Anne-Marie
Slaughter

SYSTEM

ERROR

Where Big
Tech Went
Wrong and
How We
Can **Reboot**

Rob Reich
Mehran Sahami
Jeremy
M. Weinstein





Stanford University
Human-Centered
Artificial Intelligence

Search this site



About ▾ Centers ▾ Research ▾ Education ▾ Policy ▾ News ▾ Events ▾

Advancing AI research, education, and policy to improve the human condition



Sheng Wang | Generative AI for Multimodal
Biomedicine

On the Opportunities and Risks of Foundation Models

Rishi Bommasani* Drew A. Hudson Ehsan Adeli Russ Altman Simran Arora
Sydney von Arx Michael S. Bernstein Jeannette Bohg Antoine Bosselut Emma Brunskill
Erik Brynjolfsson Shyamal Buch Dallas Card Rodrigo Castellon Niladri Chatterji
Annie Chen Kathleen Creel Jared Quincy Davis Dorottya Demszky Chris Donahue
Moussa Doumbouya Esin Durmus Stefano Ermon John Etchemendy Kawin Ethayarajh
Li Fei-Fei Chelsea Finn Trevor Gale Lauren Gillespie Karan Goel Noah Goodman
Shelby Grossman Neel Guha Tatsunori Hashimoto Peter Henderson John Hewitt
Daniel E. Ho Jenny Hong Kyle Hsu Jing Huang Thomas Icard Saahil Jain
Dan Jurafsky Pratyusha Kalluri Siddharth Karamcheti Geoff Keeling Fereshte Khani
Omar Khattab Pang Wei Koh Mark Krass Ranjay Krishna Rohith Kuditipudi
Ananya Kumar Faisal Ladhak Mina Lee Tony Lee Jure Leskovec Isabelle Levent
Xiang Lisa Li Xuechen Li Tengyu Ma Ali Malik Christopher D. Manning
Suvir Mirchandani Eric Mitchell Zanele Munyikwa Suraj Nair Avanika Narayan
Deepak Narayanan Ben Newman Allen Nie Juan Carlos Niebles Hamed Nilforoshan
Julian Nyarko Giray Ogut Laurel Orr Isabel Papadimitriou Joon Sung Park Chris Piech
Eva Portelance Christopher Potts Aditi Raghunathan Rob Reich Hongyu Ren
Frieda Rong Yusuf Roohani Camilo Ruiz Jack Ryan Christopher Re Dorsa Sadigh
Shiori Sagawa Keshav Santhanam Andy Shih Krishnan Srinivasan Alex Tamkin
Rohan Taori Armin W. Thomas Florian Tramèr Rose E. Wang William Wang Bohan Wu
Jiajun Wu Yuhuai Wu Sang Michael Xie Michihiro Yasunaga Jiaxuan You Matei Zaharia
Michael Zhang Tianyi Zhang Xikun Zhang Yuhui Zhang Lucia Zheng Kaitlyn Zhou
Percy Liang*¹

Center for Research on Foundation Models (CRFM)
Stanford Institute for Human-Centered Artificial Intelligence (HAI)
Stanford University

AI is undergoing a paradigm shift with the rise of models (e.g., BERT, DALL-E, GPT-3) that are trained on broad data at scale and are adaptable to a wide range of downstream tasks. We call these

On Leave from Stanford: 2023-2025



The California Report on Frontier AI Policy

The innovations emerging at the frontier of artificial intelligence are poised to create historic opportunities for humanity but also raise complex policy challenges. As the epicenter of global AI innovation, California has a unique opportunity to continue supporting developments in frontier AI while addressing substantial risks that could have far-reaching consequences for the state and beyond.

Transluce



Jacob Steinhardt



Sarah Schwettmann



Tiffany Tzeng



Kevin Meng



Kaiying Hou



Dami Choi



Neil Chowdhury



Daniel Johnson



Vincent Huang



Nitarshan Rajkumar



Rob Reich

Transluce is building the public tech stack for understanding and governing AI, in order to ensure reliable and safe rollouts of frontier AI systems.

transluce.org

2024-2025: Senior Advisor to the United States AI Safety Institute

NIST

Search NIST

Menu

Information Technology / Artificial intelligence

U.S. ARTIFICIAL INTELLIGENCE SAFETY INSTITUTE

Strategic Vision

Guidelines

Artificial Intelligence Safety Institute Consortium

NIST AI Engagement

AI @ NIST

The U.S. AI Safety's Institute's (US AISI) mission is to identify, measure, and mitigate the risks of advanced AI systems so that we can harness the enormous potential of this breakthrough technology. US AISI is tasked with developing the testing, evaluations, and guidelines that will help accelerate trustworthy AI innovation in the United States and around the world – with a keen focus on helping to prevent misuse of this technology by those who seek to undermine our public safety and national security.

US AISI is housed within the Commerce Department as a part of NIST and draws on NIST's time-tested scientifically grounded processes to facilitate the development of trusted standards around new technologies.

NEWS AND UPDATES



Technical Blog: Strengthening AI Agent Hijacking Evaluations

JANUARY 17, 2025

Lama AI models are increasingly used to power agentic systems, or "agents" which can automate complex tasks

2024-2025: Biden Admin (US AISI)

Bureaus and offices • Contact us

U.S. Department of Commerce

Search

Search

ES NEWS DATA AND REPORTS WORK WITH US

Home » News » Press Releases

Statement from U.S. Secretary of Commerce Howard Lutnick on Transforming the U.S. AI Safety Institute into the Pro-Innovation, Pro-Science U.S. Center for AI Standards and Innovation

2024-2025: Trump Admin (US CAISI)

AI Governance Lab at Stanford University

I created the AI Governance Lab at Stanford University to bring together philosophers, social scientists and policy experts, and technologists with the aim of shaping the AI transition. What AI will do to our personal, social, economic, and political lives is a design question, not a prediction problem. And creating the technical, social, and political arrangements to shape AI's impact on the world is one of the most important problems of our age.

I am Rob Reich, a professor of political science at Stanford University, with courtesy appointments in Philosophy, Law, and Education. I teach and write about the governance of frontier science and technology. In the past, I've written about philanthropy in democratic societies, and the purposes of public education.

I serve as associate director of the Institute for Human-Centered AI and as faculty director of the Ethics, Society, and Technology Initiatives at the Center for Ethics in Society.

Why subscribe?

Subscribe to get full access to the newsletter, early invitations to special webinars and

A Note on Course Policy





Devices in the Classroom

Research on tech use in the classroom shows:

1. Writing by hand increases your retention and engagement
2. Using a laptop in classes decreases learning
 - standing temptation to multi-task
 - even if not distracted, worse conceptual learning as compared to handwriting
3. Distract others around you. Second-hand smoke effect

Handwriting notes >
Typing notes >
AI summaries



Handwriting Increases Retention

Research Article

The Pen Is Mightier Than the Keyboard: Advantages of Longhand Over Laptop Note Taking



Pam A. Mueller¹ and Daniel M. Oppenheimer²

¹Princeton University and ²University of California, Los Angeles

Handwriting but not typewriting leads to widespread brain connectivity: a high-density EEG study with implications for the classroom



F. R. (Ruud) Van der Weel



Audrey L. H. Van der Meer*

Developmental Neuroscience Laboratory, Department of Psychology, Norwegian University of Science and Technology, Trondheim, Norway

Over a decade of research supports the idea that writing by hand increases retention... New brain data supports this!

Why? “[W]hen taking notes by hand, it’s often impossible to write everything down; **students have to actively pay attention** to the incoming information and process it—prioritize it, consolidate it and try to relate it to things they’ve learned before. This conscious action of building onto existing knowledge can make it easier to stay engaged and grasp new concepts.” ([SciAm, 2024](#))

Digital Distraction

Digital distraction is when your focus shifts from an academic to a non-academic task. (Reducing focus and increasing cognitive load.)

Research shows how devices enable this: think about how many times you check your email, texts, social media, or browse the internet or texts during homework or class.

And it impacts those nearby! Research shows that regardless of where laptops are near the student, classroom laptop use negatively impacts learning. ([Hall, 2020](#))

If you're curious, track how many times you switch between academic and non-academic tasks during lecture or homework.





The Secondhand Smoke Effect

How nearby device use harms *your* learning

Perhaps the most compelling evidence concerns students who aren't using devices but sit near those who are.

Large Lectures Amplify Effects

1

Unavoidable Secondhand Effects

In a 100-person lecture hall, students cannot choose to sit away from all device users. The distraction externality becomes unavoidable.

2

Longer Presentations

Research shows handwriting advantages are more pronounced for pronounced for longer presentations.

3

Passive Learning Risk

Large lectures already tend toward passive reception. Active processing forced by handwriting counteracts this inherent inherent passivity.


4

Impossible Monitoring

Instructors cannot observe off-task behavior in large classes, classes, making policy-based approaches more practical than case-by-case intervention.

Crown 290: Tech-Free Zone

- No laptops, tablets, or AI note takers. [But dedicated, no internet digital notetaking via stylus is allowed; eg., ReMarkable tablet]
- And remember that we post all slides online, so no need to worry about catching every detail during class
- Learn the skill of taking notes



“We have Paleolithic emotions, medieval institutions, and godlike technology.”

– E.O. Wilson

What do we mean by AI?



Algorithms

Written decision procedures. The oldest and broadest sense of the term.



Machine learning and neural networks

Systems that infer patterns from data instead of following rules a person wrote down.



Generative AI (LLMs)

Models trained to produce text, images, audio and code in response to a prompt.



AI agents

Models given tools, memory and autonomy to carry out multi-step tasks with little supervision.



Embodied AI and AGI

Robotics and general-purpose capability. Largely a projection about where the technology is heading.



World models

Systems that build an internal simulation of an environment and predict how it will behave.

Four Schools of Thought on AI



The AI Doomers

AI Boomers

AI As Normal Technology

The “Tragic Accelerators”

THE AI DOOMERS

Advanced AI could escape human control or concentrate dangerous power.

	Geoffrey Hinton “I don’t think they should scale this up more until they have understood whether they can control it.” The New York Times, May 2023	Mitigating the risk of extinction from AI should be a global priority alongside other societal-scale risks such as pandemics and nuclear war. Signatories: ☒ AI Scientists ☒ Other Notable Figures Geoffrey Hinton Emeritus Professor of Computer Science, University of Toronto Yoshua Bengio Professor of Computer Science, U. Montreal / Mila Demis Hassabis CEO, Google DeepMind Sam Altman CEO, OpenAI Dario Amodei CEO, Anthropic Dawn Song Professor of Computer Science, UC Berkeley Ted Lieu Congressman, US House of Representatives Bill Gates
	Yoshua Bengio “Autonomous superhuman AI systems have the potential to be superdangerous, in ways that previous technologies weren’t.” FAQ on Catastrophic AI Risks, 2023	
	Stuart Russell “How do we maintain power, forever, over entities that will eventually become more powerful than us?” Senate testimony, July 2023	
	Meredith Whittaker “...these are technologies controlled by a handful of corporations who will ultimately make the decisions about what technologies are made...” Fast Company interview, May 2023	

AI BOOMERS

Rapid development promises major benefits; excessive restraint could forfeit them.



David Sacks

“I hope that people will become a little bit more optimistic about this industry as more and more miraculous products come out.”
[Davos conversation, Jan. 2026](#)



Marc Andreessen

“AI will not destroy the world, and in fact may save it.”
[Why AI Will Save the World, June 2023](#)



Jensen Huang

“Every industry, every company, every country must produce a new industrial revolution.”
[NVIDIA AI Summit Japan, Nov. 2024](#)



Mark Zuckerberg

“I am extremely optimistic that superintelligence will help humanity accelerate our pace of progress.”
[Personal Superintelligence, July 2025](#)



Donald Trump

““I’m not going to stifle growth of something that will be bigger than the industrial revolution.” ... “Whoever wins superintelligence, wins”
[U.N. General Assembly, Sept. 2026](#)



The only control or “guardrails” that AI needs is a STRONG AND SMART (High IQ!) PRESIDENT, and the U.S.A. has that, in spades! The Trump Administration has stopped AI “people” from doing bad, or potentially bad, “things,” like Dario (Anthropic!), who is now pretending to be a “perfect little angel” - and we will continue to do so! We already have tremendous CRIMINAL and REGULATORY power over these companies! There is a SICK conspiracy going on against AI and Data Centers, and the only one that is happy about it is China. WHOEVER WINS AI, WINS! We are leading China, and all others, and will continue to do so. Conspiracy Theorists, Treasonists, Traitors, and Leakers, BEWARE! Thank you for your attention to this matter! President DONALD J. TRUMP

AI AS NORMAL TECHNOLOGY

Powerful technology that people and institutions can shape and govern.



Arvind Narayanan

“We view AI as a tool that we can and should remain in control of.”

[With Sayash Kapoor, AI as Normal Technology, 2025](#)



Sayash Kapoor

“Even transformative, general-purpose technologies such as electricity and the internet are ‘normal’ in our conception.”

[With Arvind Narayanan, AI as Normal Technology, 2025](#)



Gary Marcus

“...deep learning was only a tiny part of what we need to build if we’re ever going to get trustworthy AI.”

[Deep Learning Is Hitting a Wall, 2022](#)



Yann LeCun

“A baby learns how the world works in the first few months of life. We don’t know how to do this [with AI].”

[TIME interview, Feb. 2024](#)



Fei-Fei Li

“There’s nothing ‘artificial’ about AI. It’s inspired by people, it’s created by people, and—most importantly—it impacts people.”

[House Science Committee testimony, 2018](#)

THE “TRAGIC ACCELERATORS”

AI promises major benefits, but progress may outrun the safeguards needed to manage its risks.



Dario Amodio

“This recursive self-improvement could outrun our ability to understand and control these systems, and so must be pursued very carefully, if at all.”

[We Must Pace the Frontier, Sept. 2026](#)



Sam Altman

“AI will probably most likely lead to the end of the world, but in the meantime, there’ll be great companies.”


[Open Air fireside chat, 2015](#)



Elon Musk

“I think there’s about a 20% chance that AI leads to human extinction.”

[Joe Rogan interview, 2025](#)



Demis Hassabis

“When it comes to very powerful technologies—and obviously AI is going to be one of the most powerful ever—we need to be careful.”

[TIME interview, Jan. 2023](#)

The AI Stack: Different layers, Different Governance



Applications

User-facing products and autonomous agents

Chatbots · copilots · coding agents · enterprise tools



Models

Pre-trained foundation models, fine-tuning and APIs

GPT · Claude · Gemini · Llama · Mistral · DeepSeek



Data

Training corpora, curation, labeling and evaluation

Web & code · licensed · human-labeled · synthetic



Chips & Infrastructure

Specialized silicon, data centers and cloud platforms

Nvidia GPUs · TPUs · TSMC fabs · AWS · Azure · Google Cloud



Energy

Power generation, grid capacity, transmission and cooling

Fossil fuels · nuclear · renewables · water

DEFINITIONS

What do we mean by “governance”?

REGULATION

Public rules, public enforcement

National law	Federal, state and local legislation
Executive action	Orders, procurement rules, export controls
Courts	Judicial decisions and the common law
Other jurisdictions	EU AI Act, China, Korea, UK
International	Treaties and multilateral bodies

EXTRALEGAL SOURCES


Private rules, private enforcement

Insurers	Coverage conditions and exclusions
Investors	Warranties and diligence
Customers	Procurement terms and audit rights
Platforms	Acceptable-use policies, access
Auditors & standards	Certification and benchmarks

Also on the map: corporate boards, lawyers, and professional associations — internal governance that answers to both columns.

The AI Accountability Ecosystem

Each actor applies its own standards — ISO, NIST, NVCA, and more — to condition what AI firms can build, ship, and sell.



Unique Challenges to AI Governance

◆ THE TECHNOLOGY RESISTS CONTROL



Speed of development and ubiquity of applications

Capability advances on compute's timescale; law moves on the timescale of institutions. The pacing gap widens each cycle.



Stochastic, not deterministic

Behavior can't be fully specified, tested, or guaranteed — the opposite of the compliance-by-rule model regulation assumes.



Agentic

Systems increasingly act in the world, not just generate content. Delegation, attribution, and liability all get harder.



Escalating capability (AGI)

A general, moving target that may eventually outrun the capacity of those governing it. You regulate a horizon, not a fixed artifact.

◆ THE ECOSYSTEM RESISTS CONTROL



Geopolitical race

Unilateral restraint cedes ground to rivals, and no single jurisdiction can govern a global technology.



Open weight models

Once weights are public they can't be recalled, and their guardrails can be stripped. Intervention becomes irreversible.



Governing across the stack

Compute, models, and applications each offer different leverage — where to regulate is as contested as whether.



Talent Gap

The effectiveness of any regulatory regime depends on adequate enforcement resources and talent in government.

The recurring design choice: Regulate the technology, or regulate its applications?

A Taxonomy of AI Risks



Malicious use & weaponization

intentional harm by bad actors

- ▶ CBRN uplift — chemical, biological, radiological, nuclear
- ▶ Offensive cyber & critical-infrastructure attacks
- ▶ Disinformation & influence ops; deepfakes
- ▶ Fraud, scams & non-consensual imagery (incl. CSAM)
- ▶ Autonomous weapons; mass surveillance & repression



Malfunction & loss of control

unintended technical failure

- ▶ Loss of control of autonomous agents; power-seeking
- ▶ Unreliable or inaccurate outputs in critical fields
- ▶ Multi-agent failures — flash crashes, collusion
- ▶ Security holes: poisoning, model theft, prompt injection



Societal, economic & environmental

structural & systemic

- ▶ Labor disruption & job displacement
- ▶ Power & wealth concentration
- ▶ Erosion of democracy & information integrity
- ▶ Environmental & climate costs — energy, water, carbon
- ▶ Intellectual-property & creative-economy harms



Individual, psychological & rights

person-level harm

- ▶ Psychological & mental-health harms; over-reliance
- ▶ Bias, discrimination & disenfranchisement
- ▶ Privacy violations & individual surveillance
- ▶ Child safety & healthy development
- ▶ Human dignity, autonomy & human rights

Every risk we catalogued is a risk of AI being used.

There is a parallel risk of AI not being used – or being used only by others.

Outside of the countries developing foundation models, the cost of slow diffusion may exceed the cost of governance failures.

Imperfect Regulatory Analogies

ANALOGY → implied model	WHY IT FITS	WHY IT MISLEADS
 Nuclear → <i>international regime; input controls</i>	✓ Catastrophic and dual-use; control a scarce input via an international regime.	✗ Weights are copyable software, not countable fissile material.
 Electricity → <i>utility rules; point-of-use safety</i>	✓ A general-purpose technology you regulate at the point of use.	✗ Inert and predictable — no agency, emergence, or rivalry.
 The internet → <i>multistakeholder standards</i>	✓ Global, layered, permissionless; a familiar jurisdiction problem.	✗ Open plumbing no one owns — frontier AI is concentrated.
 Social media → <i>transparency, audits, liability</i>	✓ Same firms and network effects; the audit toolkit transfers.	✗ Platforms move content; AI generates it and acts.
 Recombinant DNA (Asilomar) → <i>self-governance; moratoria</i>	✓ Scientists pausing to set safety norms for dual-use research.	✗ A small, legible field then; AI is global, commercial, and racing.
 Aviation → <i>pre-market certification</i>	✓ Certification and incident reporting build safety and trust.	✗ Airframes are fixed and deterministic; models are not.
 Pharmaceuticals (FDA) → <i>staged trials; approval</i>	✓ Phased trials and post-market surveillance for uncertain risks.	✗ One general-purpose model isn't one drug for one indication.
 Climate / IPCC → <i>assessment body; treaties</i>	✓ A shared global risk needing an assessment body and coordination.	✗ The harm here is contested and fast, not measured and gradual.

Common AI Governance Requirements

The shared playbook the private accountability ecosystem is converging on

1

Predeployment testing

Red-teaming, evals, and benchmarking before release

2

Transparency & documentation

Model cards, data sheets, disclosure to counterparties

3

Incident reporting & monitoring

Logging, notification duties, post-deployment watch

4

Accountability & governance

A named owner and board-level oversight of AI risk

5

Risk-management process

A repeatable, framework-mapped assessment cycle

6

AI system inventory

A catalog of deployed systems and their risk tiers

7

Data governance & provenance

Training-data lineage, IP rights, privacy compliance

8

Human oversight & control

Escalation paths, rollback, and kill-switch capability

9

Third-party / supply-chain risk

Governing vendor and foundation-model dependencies

10

Security controls

Access management, robustness, misuse prevention

11

Auditability & evidence

Audit trails and attestations a third party can inspect

The Law of AI

I

Constitutional

II

Statutory

III

Common law

AI and the Constitution



RIGHTS AGAINST THE STATE

● First Amendment

Is model output “speech,” and whose? Labeling mandates, deepfake bans, code as speech.

● Fourth Amendment

Surveillance and inference: face recognition, predictive policing, data aggregation.

● Due process

Automated state decisions: notice, explanation, and the right to contest a score.

● Equal protection

Algorithmic bias: proxies, the intent requirement, and race-conscious fixes.

STRUCTURE OF GOVERNMENT

● Separation of powers

Major questions, nondelegation, unitary executive; can agencies delegate to AI?

● Federalism

State AI laws vs. federal preemption; dormant Commerce Clause and extraterritoriality.

● War & foreign affairs

Autonomous weapons, war powers, and export controls on chips and models.

● Judicial power

Standing for algorithmic harm; AI-assisted judging; the jury’s role.

DEMOCRACY, JUSTICE & LIBERTY

● Elections & democracy

Political speech and the Elections Clause: deepfakes, microtargeting, campaign AI.

● Criminal procedure

Confrontation, counsel, and self-incrimination: machine evidence, source-code access.

● Personhood & authorship

Rights for AI? “Authors” and “Inventors” in the IP Clause; personhood analogies.

● Privacy & autonomy

Informational privacy and intimate autonomy amid AI companions and inference.

AI and Statutory Law

Statutes assume human actors, human-scale conduct, and finished products. AI is none of these.



INFORMATION & GOVERNMENT

● Platforms & speech

Does §230 cover generated content? Deepfake, intimate-image, likeness, and robocall laws.

● Intellectual property

Training and fair use; human authorship and inventorship; trade secrets.

● Privacy & data

Consent can't govern inference; sectoral and state laws vs. training on personal data.

● Elections & government

Deepfake rules for campaign ads; APA, FOIA, and procurement rules for government AI.

MARKETS & LIVELIHOODS

● Competition

Algorithmic collusion without agreement; concentration across the AI stack.

● Consumer protection

Deception, manipulation, and "AI-washing"; companion bots and kids' online safety.

● Finance & insurance

Credit and underwriting explanations, securities disclosure, and robo-advice.

● Civil rights & work

Disparate impact in hiring, housing, and credit; algorithmic management of workers.

HEALTH, SAFETY & SECURITY

● Health & professions

FDA review of models that keep learning; licensing and unauthorized-practice laws.

● Transport & workplace safety

Rules that assume a human driver or pilot; OSHA and autonomous equipment.

● Environment & energy

Data-center permitting under NEPA and the Clean Air Act; grid, emissions, and water.

● Crime & national security

Fraud, hacking, and deepfake-porn statutes; export controls and investment screening.

FOUR QUESTIONS FOR EVERY STATUTE

Definitions

Is AI a "person" or an "author"?

Duty-bearer

Developer, deployer, or user?

Federalism

State patchwork vs. preemption

Enforcement

Agency capacity; who may sue

AI and the Common Law

Doctrines built for human actors meet a defendant that is opaque, autonomous, and everywhere.



TORTS

● Negligence & products liability

Reasonable care for developers and deployers; is a model a “product”; causation without explanation.

● Defamation & dignitary harms

Hallucinated falsehoods, deepfakes, cloned voices and likenesses: who is the “speaker”? No §230 shield.

● Duty & novel harms

Duties to at-risk users and third parties; scraping as trespass to chattels; data centers as nuisance.

AGENCY, CONTRACT & DUTY

● Agency

Can an AI agent bind its principal? Actual vs. apparent authority; vicarious liability for its torts.

● Contract

Formation by and between agents; assent and mistake; terms of service, waivers, and unconscionability.

● Professional duties

Fiduciary and malpractice standards when professionals use AI, or fail to; loyalty owed by AI advisers.

PROPERTY, MENS REA & PROOF

● Property & unjust enrichment

Training on others’ data and likeness: conversion, misappropriation, restitution. Who owns outputs?

● Mental states & personhood

Intent, knowledge, and foreseeability for a non-human actor: must liability always trace back to a person?


● Evidence & procedure

Authenticating deepfakes, admitting AI-generated proof, discovering models and chats, hallucinated citations.

STANFORD
LAW SCHOOL | **AI INITIATIVE**

<https://law.stanford.edu/ai-initiative/>





Anka Reuel

- Computer Science PhD Candidate @ Stanford
- Incoming Assistant Prof @ Harvard (w/ appointments in Public Policy & CS)
- Former Vice Chair for GPAI Code of Practice Process @ European Commission
- Former member of expert working group convened by Gov. Newsom to inform SB-53
- Things I enjoy: Kung Fu Panda, playing PS5/Switch, traveling, triathlon



What I Want You To Take Away From This Class

What I Want You To Take Away From This Class

1) “AI” is usually too broad a category for good governance.

If an intervention is sensible depends on the system, use, actor, capability, and risk you are actually trying to govern.

What I Want You To Take Away From This Class

- 1) “AI” is usually too broad a category for good governance.
- 2) Governance needs measurements that are meaningful, not just available.

Governance often relies on proxies. Those proxies can be imperfect, gameable, or become outdated as the technology changes.

What I Want You To Take Away From This Class

- 1) “AI” is usually too broad a category for good governance.
- 2) Governance needs measurements that are meaningful, not just available.
- 3) Governance interventions should not be treated as permanent.

What works today can stop working as the technology changes.

Example 1: Compute Thresholds

Regulation	Compute Threshold (in FLOPs*)	What it does
EU AI Act	$>10^{25}$	Presumes a GPAI model has high-impact capabilities and therefore systemic risk
California's SB 53	$>10^{26}$	Defines a frontier model; scopes frontier-model transparency and safety requirements
New York RAISE Act	$>10^{26}$	Defines a frontier model; scopes frontier-model transparency and safety requirements
South Korea AI Basic Act	$>10^{26}$	Helps define AI systems subject to additional safety obligations

*FLOP: A floating point operation is a single arithmetic calculation; total training compute counts every FLOP performed across a model training run

Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...


Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...

Example 1: Compute Thresholds



(a) MMLU-Pro (knowledge): weaker saturation. (b) MATH Lvl 5 (reasoning): stronger saturation.

Figure 6. Task-dependent saturation on Open LLM Leaderboard v2. “Dominated large” marks large (> 13B) models whose task score is below the cumulative best score achieved by small models up to that date (the small-model boundary). MMLU-Pro appears less saturated than MATH Lvl 5.

Orange: newer/post-cutoff small models <13B from PROTEUS-2K ; Gray hollow: large models whose performance is already beaten by a small model available by that date ;
Dashed line: March 13, 2025, when Open LLM Leaderboard v2 retired

Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...


Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...

Example 1: Compute Thresholds



Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...

Example 1: Compute Thresholds

Crude proxy for capability (& risk) (if it is at all a reliable proxy!)

Unclear what risks governments are concerned about

Underspecified: inference time capability increases, agents, ...

What I Want You To Take Away From This Class

- 1) “AI” is usually too broad a category for good governance.
- 2) Governance needs measurements that are meaningful, not just available.
- 3) Governance interventions should not be treated as permanent.
- 4) AI governance is not the same thing as AI regulation.


See Nate's slides.

What I Want You To Take Away From This Class

- 1) “AI” is usually too broad a category for good governance.
- 2) Governance needs measurements that are meaningful, not just available.
- 3) Governance interventions should not be treated as permanent.
- 4) AI governance is not the same thing as AI regulation.
- 5) Every governance intervention is a bet on how the world works.

It embeds assumptions about the technology, human behavior, institutions, and incentives.

Example 2: Human Oversight Over Agents




Example 2: Human Oversight Over Agents

Keeping humans in control while enabling agent autonomy

A central tension in agent design is balancing agent autonomy with human oversight. Agents must be able to work autonomously—their independent operation is exactly what makes them valuable. But humans should retain control over how their goals are pursued, particularly before high-stakes decisions are made. For example, an agent helping with expense management might identify that the company is overspending on software subscriptions. Before it starts cancelling subscriptions or downgrading service tiers, the company would likely want a human to give approval.

In Claude Code, humans can stop Claude whenever they want and redirect its approach. It has read-only permissions by default, meaning it can analyze and review information within the directory it's initialized in without asking for human approval, but must ask for human approval before taking any actions that modify code or systems. Users can grant persistent permissions for routine tasks they trust Claude to handle.

Example 2: Human Oversight Over Agents




Example 2: Human Oversight Over Agents

"Claude Code users approve 93% of permission prompts. [...]

Over time that leads to approval fatigue, where people stop paying close attention to what they're approving."

Example 2: Human Oversight Over Agents



Example 2: Human Oversight Over Agents

AI • CODING

An AI-powered coding tool wiped out a software company's database, then apologized for a 'catastrophic failure on my part'

By **Beatrice Nolan**
Tech Reporter

Follow



July 23, 2025, 7:22 AM ET

What I Want You To Take Away From This Class

- 1) “AI” is usually too broad a category for good governance.
- 2) Governance needs measurements that are meaningful, not just available.
- 3) Governance interventions should not be treated as permanent.
- 4) AI governance is not the same thing as AI regulation.
- 5) Every governance intervention is a bet on how the world works.

A few more thoughts for the class

1. We don't want to tell you what to think.
2. This class is meant to be a conversation.
3. AI governance needs many voices, including yours.



AI ACTION SUMMIT

Logistics



Grading

- AI Evaluation Research Project – 60%
 - 4-page question, relevance, and novelty write-up – 25%
 - Project Plan – 20%
 - Peer Review – 10%
 - Final Paper – 45%
- Public Comment Assignment – 30%
- Attendance/Participation – 10%

Public Commentary

- Respond to the Request for Comments (RfC) on the [Colorado Automated Decision-Making Technology and Chatbot Safety Acts](#)
- 1000 to 1500 words
- Deadline for commentary in this class: Oct 23rd @ 5pm PST
- Actual RfC deadline: Oct 26
- You don't have to actually submit your comment but we strongly encourage it (it creates a public artifact for you!)

AI Evaluation Research Project

- Track 1: Who decides what gets measured?
 - Map who funds, builds, and uses evaluations
 - Identify gaps, incentives, and conflicts of interest
 - Design institutions or mechanisms for setting evaluation priorities
- Track 2: Building evaluations that deserve trust
 - Test reliability or validity of an existing evaluation
 - Quantify uncertainty or study a measurement failure mode
 - Design better reporting or documentation standards
- Track 3: Evaluations at the point of decision
 - Study how evaluations are used in regulation, procurement, litigation, or deployment
 - Propose evidence standards or decision thresholds for specific use cases
 - Analyze what happens when evidence is ambiguous or contested

AI Evaluation Research Project

- Milestone 1: Question, relevance, and novelty write-up – due Oct 16th at 5PM
- Milestone 2: Project Plan – due October 30th at 5PM
- Milestone 3: Peer Review – due Nov 6th at 5PM
- Milestone 4: Final Paper – due Dec 4th at 5PM
- You can work alone or with one team member

Attendance

- You get two free missed classes.
- Attendance will be taken at the beginning of each class.
- Any student coming in late shall be marked as absent.
- We cannot make exceptions to these rules given the number of people in this class

Schedule

- Week 2: Technical Foundations (pre-LLM era, LLMs, agents, ...)
- Week 3: AI Evaluations; Bias, Discrimination, and Fairness
- Week 4: Copyright, Creativity, & Privacy; Defamation & Mental Health
- Week 5: AI & Democracy; AI Infrastructure & the Environment
- Week 6: Existential Risks & Benefits, AGI, Autonomous R&D, CBRN; EU Law and Policy

Schedule

- Week 7: Federal & State-Level US Law and Policy; Rest of the World Approaches to AI Governance
- Week 8: Geopolitics and Global AI Governance; Competition, Open Models, & Antitrust
- Week 9: The Future of Labor; What Does It Mean to Be Human in an AI Age?
- Week 10: Outlook & AI Futures

Course Communications

- Check the syllabus & Ed first and see if you find your answer there (especially if it's about logistics).
- If you have any remaining questions regarding the class or need to miss a lecture for a valid reason, please email any of the TAs first
- Exceptions: Any sensitive information you're not comfortable sharing with the TAs; in that case, email any of the instructors
- If your question or comment is likely shared by other students, please use Ed so everyone can benefit from the answer

Team

Course staff



Dominic Zappia

Head Course Assistant

zappia@stanford.edu



Roberta Fischli

Co-Head Course Assistant

fischli@stanford.edu



Yifei Cheng

Course Assistant

yifechen@stanford.edu



Bruria Friedman-Feldman

Course Assistant (SLS)

bruriaff@stanford.edu



Thay Graciano

Course Assistant

thayg@stanford.edu



Yasi Khan

Course Assistant

yrkhan@stanford.edu



Tia Vasudeva

Course Assistant

tvasudev@stanford.edu

Please add technical terms & questions you want me to
discuss

