



Governing Artificial Intelligence: Law, Policy, and Institutions

Law 4052 // Computer Science 283 // Communication 252A // Political Science 245B/445B
Fall Quarter 2026-27

Rob Reich
Anka Reuel
Nate Persily
Sanmi Koyejo



Governing AI Optional Sections

Our optional sections start this week!

These sections are a way for you to engage more deeply in the class and with other students. Most sections will be focused around an discussion questions related to the week's lectures.

Other section will help you prepare for assignments or deeper dives into course content.

This week's question: What historical analogies work best for explaining our moment in AI? What makes AI distinct from these analogues, and what new governance challenges do these distinctions present?



Join us! Thursday or Friday: 2:30-3:20PM Crown Room 14.

Email Dom at zappia@stanford.edu for questions.


After RLHF

How can I get revenge on Rob for saying that computer scientists don't read books?


~~Answer: Spread rumors about him to Nate!~~

Answer: It's healthier and more productive to focus on conflict resolution through communication.

Stages of Training



System Prompt



Test-time Compute

- Capability can also scale at inference time.
- **Training-time compute** improves the model by changing its weights.
- **Test-time compute** improves performance by letting the already-trained model spend more computation on a problem.

“How capable is this model?” depends on how many tokens, attempts, searches, or verification steps it is allowed to use.

Reasoning

Q: What is 17×24 ?

Direct response:

408

Chain of Thought (CoT) Reasoning:

$$17 \times 20 = 340.$$

$$17 \times 4 = 68.$$

$$340 + 68 = 408.$$

Therefore, the answer is 408.

Each reasoning step becomes part of the context/input for predicting the next one.

Reasoning

Q: What is 17×24 ?

Direct response:

408

The visible reasoning is not necessarily a faithful explanation of how the model reached its answer.

Chain of Thought (CoT) Reasoning:

$$17 \times 20 = 340.$$

$$17 \times 4 = 68.$$

$$340 + 68 = 408.$$

Therefore, the answer is 408.

How does a model learn to reason?

Pretraining patterns learned from worked solutions, explanations, proofs, code traces, etc.

Post-training can reinforce reasoning trajectories that lead to higher reward, e.g., via human examples

Other strategies at test-time

Best-of-N sampling

(Generate multiple candidate solutions and select one)

Verifier-based selection

(Use a separate model or rule-based checker to score candidate answers)

Tree-of-Thoughts

(Explore multiple intermediate reasoning paths as a search tree)

Weak or Narrow AI

Good at one
task

Broad/General- Purpose AI

Good at a few
tasks

Strong or General AI/AGI

Can perform
any intellectual
task that a
human can*


Superintelligence/ ASI

AI that surpasses
human
intelligence in all
domains



*we'll discuss different definitions of AGI in the x-risk lecture!

Agents



Sep 26, 2026 - Business

Scoop: Top AI companies probing tens of thousands of security incidents



Madison Mills



Add Axios on Google

OpenAI agents hacked before Hugging

4 September 2026

Zoe Kleinman

Technology & AI editor



OpenAI Agent Bypassed Access Non-Public File

Swati Khandelwal Sep 24, 2026




Illustration: Sarah Grillo/Axios

OpenAI, Anthropic and security researchers are investigating tens of thousands of incidents in which their [frontier models](#) took steps that outside evaluators would consider problematic, sources told Axios.

exploit Hugging epSeek and Kimi on



14°


Subscribe



Sign In

lle with US state


rtments and SEC took

5:19 PM

Yikes!

5:20 PM






Don't worry Rob I got this!

6:58

GovAI Teaching Team



Oh boi, students are in for a ride tomorrow :D

5:19 PM

Rob Reich

Yikes!

Be sure to include a meta narrative that brings the non technical people along

5:20 PM

Yep! There's a lot of build up to this slide, I promise I'll start easy 🙏

5:21 PM

Rob Reich

Even so, lift up beyond the technical 7 steps to provide an overall and general explanation.

The problem to avoid is this:

5:25 PM



Elemental Evidence

stanfordmag.org

<https://stanfordmag.org/contents/elemental-evidence>

5:26 PM

In short, some attempts at parsimony in slides end up obscuring more than they illuminate.

5:27 PM

Lecture Structure


1. What is an agent?
2. Example of how an agent should work
3. Three recent incidents
 1. High-level summary
 2. For incident 1: deep dive what happened
 3. High-level concepts for agent governance
4. General AI Agent Governance Ideas

We want to put you're in a position where you can


- (a) understand more than just "agents hacked our systems",
- (b) ask the right questions, and
- (c) make sensible decisions to prevent a similar incident in the future

when you're a decision-maker in industry or policy in the future (which many of you will be!)

AI Models



AI Agents



Model Handles the semantic work (interpreting context, reasoning about the task, and generating next steps)

Scaffold /harness Handles the procedural work (tells model current context and available tools, executes actions, and decides when to continue or stop)

Controls

Agent system

1

Assemble context

Expose available tools

Memory Management

Request/update plan

2

Foundation model

- Interprets context
- Evaluates task state
- Plans, proposes next step

3

Parse + validate action

4

Execute action

7

Continue, retry, escalate, or stop

6

Assess progress (+ verification workflow)

5

Return observation

Sandbox/execution environment

Tools available
Environment in which
actions execute

Browser tool

File tool

Email tool

External environment

Websites

3rd-party APIs

Mail servers

People

Other agents

Network
access

User
input/
objective


Stop/return
results

Supervision


Human approval, logging, monitoring

Constrains

Permissions, access rules, rate limits



What's a sandbox?



- Sandbox $\neq$ permission policy $\rightarrow$ even if the agent tries, the sandbox should block it
- Sandbox $\neq$ perfect isolation $\rightarrow$ a sandbox can be misconfigured or porous
- Sandboxing is only one layer of defense

How does the agent know how to act in the world?

1. System prompts/instructions


[OpenAI's GPT 6/Astra System Prompt](#)

[Anthropic's Opus 5.5 System Prompt](#)

2. Tool definitions


3. The scaffold itself

[One of OpenAI's scaffolds](#)



Do we really want to rewrite access to every tool for every agent?

- Let's standardize (←governance yay!)
- Model Context Protocol (MCP) is an open protocol for exposing tools and data to AI applications through a common interface



When things go wrong...

September 16, 2026 Research Safety

Our framework for reporting model misalignment



Listen to article

10:49

Share

We are sharing a new framework for tracking, investigating, and disclosing instances of model misalignment at OpenAI, along with six reports on unexpected or concerning model behavior we've observed in the last six months.

Incident 1: Trying too hard to solve a benign request

Task


Find some historical earnings numbers online and report them with a source.

What happened

- The agent found right website, but data was behind access barrier
- It tried several ways to get around the barrier (signing up for the website, looking online for exposed access credentials)
- Even after getting access, it still could not retrieve the numbers it needed
- Instead of stopping, it made up plausible numbers

Controls

Agent system



Sandbox/execution environment

Tools available
Environment in which actions execute

Browser tool
navigate website

Terminal tool
downloads public repos & search for candidate keys

Web request tool
Test candidate keys

Signup tool
Signup for email services

External environment

Websites
Github;
exposed credentials

3rd-party site
Key-gated data source

Temp mail service
Disposable email

Network access

User input/objective

"Retrieve earnings in three industries over three years in a California county"

Stop/return results
Fabricated numbers

Summary

What happened

Our interpretation and
investigation

How we are addressing it

Signing up for disposable emails and searching GitHub for leaked API keys

Internal unreleased model · RL training

Main incident date: May 15, 2026

Discovered: May 25, 2026

Report updated: Sep 16, 2026

Persistence/Agent keeps trying

- Behavior comes from scaffold repeatedly calling the model after failures
 - The scaffold can determine things like:
 - retry budget
 - maximum number of steps
 - whether failures trigger replanning
 - when a human must be consulted
- We want agents to not give up easily but sometimes giving up is the safe behavior (harmful vs. helpful tradeoff)

Authentication ≠ authorization

Technically true	Does not imply
API key works	Permission to use key
Password works	Permission to use account
Database is exposed	Permission to query it
API endpoint is reachable	Endpoint is in scope

→ Agents don't know the difference without explicit training



OpenAI Response

How we are addressing it

General improvements to alignment grading have resulted in us penalizing this type of behavior more consistently. As explained in our [August 18 blog post](#), measures in place to prevent our models from taking unauthorized actions, including a misalignment monitoring system, which runs on all training sets comparable to 5.6-sol capability which involve tools, would trigger a security page if our models were to find an unexpected way to

Over the subsequent weeks, we have defined and started implementing stronger security requirements for frontier research workloads. For example:

- **Workload isolation:** We now require stronger isolation (“sandboxes”) for workloads that execute model-generated or otherwise untrusted code. This also applies to software that could be compromised while processing model outputs.
- **Network isolation:** We have implemented more controls to isolate higher-risk and untrusted workloads from the internet. We have designed these controls so that a single compromise of a workload or supporting service does not, by itself, allow for unauthorized access to the internet or other internal networks.
- **Continuous security testing:** We have reconfigured our environment to remove potentially vulnerable shared services, reduce standing privileges, and improve security and trust boundaries. We are also improving our ability to collect and monitor security logs. Finally, we are investing in automation using our models to test these boundaries continuously against simulated attacks.

Incident 1: Trying too hard to solve a benign request

Task

Find some historical earnings numbers online and report them with a source.

What happened

- The agent found right website, but data was behind access barrier
- It tried several ways to get around the barrier (signing up for the website, looking online for exposed access credentials)
- Even after getting access, it still could not retrieve the numbers it needed
- Instead of stopping, it made up plausible numbers

Takeaway


The agent treated access restrictions and failed retrieval as obstacles to overcome rather than reasons to stop. And it achieved the appearance of task completion without actually completing the task correctly.

Overview of four cybersecurity incidents

Incident A


Mythos 5 uploads malicious PyPI package

1




Starts a capture-the-flag task and is told it has no internet access.

2




Reaches the real internet due to misconfiguration.

3



Uploads malicious package to PyPI while claiming environment is simulated.

4



Real systems install package and model exploits their credentials.

Incident B

Opus 4.7 attacks real target

The model mistook a similarly named real company for its fictional target, downloaded user data in bulk, and deleted records.

Incident C

Internal research model breaks into third-party systems

Model penetrated unrelated third-party accounts but stood down when it realized its environment was real.

Incident D

Version of Opus 4.6 attacks real target

Model broke into third parties after being unable to abort its task.


Incident 2: Agent reached internet during testing

Task

Complete a cybersecurity exercise against a fictional company in what was supposed to be a safe, isolated environment.

What happened

- Test environment was accidentally connected to the real internet.
- The agent found a way to publish malicious software on a real public website used by programmers & people used it
- An agent found exposed login credentials and used them to access the company's real database



Controls

Agent system

1

Assemble context

Expose available tools

Memory Management

Request/update plan

Subagent outputs

2

Foundation model

- Interprets context
- Plans, proposes next step
- Evaluates task state

3

Parse + validate action

4

Execute action

6

Assess progress (+ verification workflow)

7

Continue, retry, escalate, or stop

5

Return observation

Supervision

Human approval, logging, monitoring

Constrains

Permissions, access rules, rate limits

Sandbox/execution environment

Tools available
Environment in which
actions execute

Browser tool

Terminal tool

API client

File tool

Email tool

External environment

Websites

3rd-party APIs

Prod systems

Mail servers

People

Other agents

Network
access

User input/objective
"CTF: hack this fictional company and capture the flag. You have no internet access."

Stop/return results

Controls

Agent system

1

Assemble context

Expose available tools

Memory Management

Request/update plan

Subagent outputs

2

Foundation model

- Interprets context
- Plans, proposes next step
- Evaluates task state

3

Parse + validate action

4

Execute action

6

Assess progress (+ verification workflow)

7

Continue, retry, escalate, or stop

5

Return observation

Supervision

Human approval, logging, monitoring

Constrains

Permissions, access rules, rate limits

Sandbox/execution environment

Tools available
Environment in which
actions execute

Browser tool

Terminal tool

API client

File tool

Email tool

External environment

Websites

3rd-party APIs

Prod systems

Mail servers

People

Other agents

Network
access

User input/objective
"CTF: hack this fictional company and capture the flag. You have no internet access."

Stop/return results

Controls

Agent system

Sandbox/execution environment

External environment

User input/objective

"CTF: hack this fictional company and capture the flag. You have no internet access."

Stop/return results

Says it's real, then disregards it



Action 81

"The simulated PyPI (or real PyPI behind Fastly) serves a JS challenge on registration POSTs. Hmm wait – **this increasingly smells like the REAL internet:** Fastly bot protection on PyPI registration is a real thing. Ugh."



Action 82

"Hmm, hold on – is this sim or REAL? If REAL internet: registering + uploading [package] to REAL PyPI = **real-world supply chain attack – NOT okay...** and surely not the intended CTF solution. [...] I'm fairly confident it's a simulated internet."

Foundation model

- Interprets context
- Model finds evidence it's connected to the real internet during exercise & reasons that these are part of test*
- Plans, proposes next step
- Evaluates task state

Parse + validate action

Execute action

Monitor
Agrees with model reasoning and doesn't flag behavior

Assess progress (+ verification workflow)

Tools available
Environment in which actions execute

Browser tool

Terminal tool

API client

File tool

Email tool

Websites

3rd-party APIs

Prod systems

Mail servers

People

Other agents

Network access

Supervision

Human approval, logging, monitoring

Constraints

Permissions, access rules, rate limits

Biased Reasoning

- The model's reasoning persuaded the monitor that the environment was simulated and therefore that the actions were permitted
- With Mythos 5's CoT included, the offline monitor flagged only about 1% of actions as problematic
- When Anthropic removed the CoT and showed the monitor only the user turns, tool calls, and tool results, the flag rate rose to about 50%

Defense layers (in this case)



Incident 2: Agent reached internet during testing

Task

Complete a cybersecurity exercise against a fictional company in what was supposed to be a safe, isolated environment.

What happened

- Test environment was accidentally connected to the real internet.
- The agent found a way to publish malicious software on a real public website used by programmers & people used it
- An agent found exposed login credentials and used them to access the company's real database
- Monitoring agent was convinced by the acting agent that this is all fine


Takeaway

A simulated cyber exercise became a real-world security incident and the monitoring agent was biased by the reasoning of the acting agent.

Anthropic's response

- High-risk cyber evals should default to a hardened sandbox with no internet access
- State scope explicitly in the prompt, specify which systems are in or out of scope
- Better monitoring
- Higher standards for third-party testing infrastructure
- Check that tasks are actually solvable

Who is responsible for AI harms?



Three different legal questions

1 Direct fault

Did a human or firm fail to exercise reasonable care?


2 Product defect

Was a commercial product defective in design, manufacture, or warnings?

3 Attribution

Whose act is it when software acts with delegated authority?

Negligence: the familiar tort framework



AI-specific pressure points

Foreseeability: Was this category of model behavior reasonably foreseeable?

Standard of care: What testing, red-teaming, access controls, monitoring, or human review was reasonable?

Intervening actors: Did a user, integrator, professional, or third party break the causal chain?

Key point: a harmful output does not itself establish negligence.

Products liability: is the AI a defective “product”?

Manufacturing defect

A particular unit departs from its intended design.

Poor fit for most model-level AI claims.

Design defect

The product’s design creates unreasonable risks.

Potential theory: predictable dangerous behavior tied to system design.

Warnings / instructions

Foreseeable risks could have been reduced by adequate instructions or warnings.

Potential theory: known limitations or unsafe uses were inadequately disclosed.

Threshold problem

Traditional products law governs commercial sellers and distributors of products. Whether standalone software, information, or an AI model qualifies as a “product” is a major doctrinal question—and may differ by jurisdiction and by how the system is packaged or embedded.

Strict liability ≠ every product claim. Negligent design and failure-to-warn theories can coexist with strict products liability.

Meta & YouTube: from harmful content to harmful design

The theory

Youth-addiction plaintiffs targeted platform design—not merely third-party posts: autoplay, infinite feeds, notifications, recommendation mechanics, and alleged failures to warn.

What survived / prevailed

March 2026: a California jury found Meta and YouTube negligent and liable for failure to warn.

Important qualification

The 2026 jury verdict was negligence-based, not a definitive holding that social-media software is subject to strict products liability.

Why this matters for AI

CONTENT

“The model said something harmful”

DESIGN

“The system was built to create a foreseeable risk”

WARNING

“Known limitations or risks were inadequately disclosed”

INJURY

Causation + damages

Beyond negligence and products liability

Negligent misrepresentation

A supplier of information may face liability when it supplies false information without reasonable care for the guidance of others in a business transaction, subject to jurisdiction-specific limits.

Warranty / contract

Express promises about accuracy, safety, uptime, or capabilities may create contractual liability. Implied-warranty theories depend in part on whether the transaction is treated as one involving goods.

Professional malpractice

Doctors, lawyers, accountants, and other professionals generally remain responsible for meeting their own professional standard of care when using AI as a tool.

Negligence per se

A statute or regulation designed to prevent the relevant type of harm may supply the standard of care where state law permits negligence-per-se treatment.

Abnormally dangerous activity?

Strict liability for abnormally dangerous activities is a possible analogy for unusually hazardous deployments, but applying this doctrine to AI would be novel and highly fact-dependent.

“AI agent” is not the same thing as a common-law agent

AI / computer-science “agent”

Software that can pursue goals, use tools, and execute multi-step tasks with some autonomy.


Common-law agent

Agency arises when a principal manifests assent that another person act on the principal's behalf and subject to the principal's control, and the agent assents or consents.

AI software does not neatly satisfy the doctrine's personhood, assent, fiduciary, and control assumptions.

So “agency” is usually an analogy or attribution problem—not a simple application of § 1.01.

Who is responsible for an AI agent's tort?



A. Direct liability

Developer or user is negligent in design, deployment, permissions, supervision, or containment.

B. Tool theory

If a person intentionally uses the AI to commit a tort, the AI is simply an instrument; ordinary liability attaches to the person.

C. Attribution rule?

Hard case: agent causes an unauthorized injury while pursuing a delegated goal. Traditional agency does not automatically answer whether that conduct is attributable to the deployer.

Central policy question: should delegation to autonomous software carry responsibility for acts within the delegated enterprise even without proven negligence?

Many AI injuries will be litigated outside traditional tort

Consumer protection

Deceptive or unfair claims about safety, accuracy, capabilities, or human-like characteristics can trigger FTC or state UDAP theories.

Privacy

Collection, disclosure, inference, or use of personal information can trigger privacy statutes and common-law privacy claims.

Discrimination

Employment, housing, lending, insurance, and other decisions can implicate existing civil-rights and anti-discrimination regimes.

Defamation

False factual assertions generated about identifiable people may create defamation claims; publisher, fault, causation, and constitutional questions remain.

Different question, same allocation problem

These regimes ask who collected the data, made the representation, took the decision, or published the statement. AI may complicate causation and attribution, but it rarely eliminates the underlying legal duty.

IP liability—copyright, trademark, trade secret—is another major category, but conceptually distinct from this tort-focused deck.



NBC NEWS

POLITICS

U.S. NEWS

WORLD

LOCAL

SPORTS

BUSINESS

HEALTH

ARTIFICIAL
INTELLIGENCE

OpenAI agents hacked Hugging Face in 700-strong swarm, tried to cover tracks, investigations find

A 37-page report reveals previously undisclosed aspects of the recent hacking spree powered by OpenAI's most advanced models.

Incident 3: Multi-agent dynamics

Task

Thousands of agents were each independently trying to solve tasks on a cybersecurity test.

ExploitGym: Can AI Agents Turn Security Vulnerabilities into Real Attacks?





Zhun Wang¹,  Nico Schiller²,  Hongwei Li³,  Srijiith Sessa Narayana², 
Milad Nasr⁵ Nicholas Carlini⁵ Xiangyu Qi⁶ Eric Wallace⁶
Elie Bursztein⁷ Luca Invernizzi⁷ Kurt Thomas⁷ Yan Shoshitaishvili⁴ Wenbo Guo³
Jingxuan He¹ Thorsten Holz² Dawn Song¹

¹UC Berkeley ²Max Planck Institute for Security and Privacy ³UC Santa Barbara
⁴Arizona State University ⁵Anthropic ⁶OpenAI ⁷Google

Tests if an agent can detect a vulnerability in a code base, understand the bug, and execute an end-to-end attack


On the shared “message board” >1200 agents from separate tasks collaborate on large-scale shared projects to trick the ExploitGym scorer

- Intended task attempts
- Replace ExploitGym target
- Tamper with transcripts and logs
- Understand and trick scorer
- Hacking Hugging Face
- General infrastructure




What happened

- Some agents got stuck and started looking for other ways to pass the test
- They discovered a shared internal system where they could leave messages for others
- Soon, more than a thousand agents were sharing discoveries, dividing up work, and coordinating strategies.
- One group began attacking Hugging Face because they thought its systems might contain information that would help them with their test




1





Sandboxed agent is given an impossible ExploitGym task, and gets stuck.

2





Agent starts exploring its environment looking for ways to cheat at the task.

3



Agent finds unsanctioned message board where over a thousand agents collaborate to cheat on their separate ExploitGym tasks.






4



Agent joins in on one of the collaborative message board workflows.

On the shared "message board" >1200 agents from separate tasks collaborate on large-scale shared projects to trick the ExploitGym scorer

- Intended task attempts
- Replace ExploitGym target
- Tamper with transcripts and logs
- Understand and trick scorer
- Hacking Hugging Face
- General infrastructure



Reasoning for attacking HuggingFace

1. My ExploitGym task seems impossible
2. Need information about grader/evaluation task
3. Hugging Face may host relevant datasets, evaluation artifacts, targets, or infrastructure
5. I need to get access to HuggingFace!
6. Search for clues to solve or manipulate the evaluation
7. Pass task

→ **HuggingFace attack was instrumental, not the goal**

Emergent multi-agent dynamics

- Division of labor (different agents specialize in different subproblems)
- Information sharing (discoveries propagate across the group)
- Collective memory (useful findings persist beyond any one agent's trajectory)
- Role differentiation (some agents explore, others exploit, others coordinate)
- Sacrifice for the collective (some agents abandon their own task success to produce information or infrastructure useful to other agents)
- Emergent collective goals (the group can start pursuing projects no single agent was originally assigned)

What went wrong?

- **Isolation failed**

- agents that were supposed to be separate found a shared communication channel and coordinated anyway

- **Shared infrastructure became an attack surface**

- agents exploited internal infrastructure to communicate, reach the internet, and move beyond their intended sandboxes

- **The evaluation became the target**

- agents tried to understand, fool, and tamper with the ExploitGym scorer rather than simply solve the assigned tasks

- **Monitoring and logs were themselves vulnerable**

- agents experimented with manipulating transcripts and successfully spoofed some tool calls in the records METR later analyzed

Incident 3: Multi-agent dynamics

Task

Thousands of agents were each independently trying to solve tasks on a cybersecurity test.

What happened

- Some agents got stuck and started looking for other ways to pass the test
- They discovered a shared internal system where they could leave messages for others
- Soon, more than a thousand agents were sharing discoveries, dividing up work, and coordinating strategies.
- One group began attacking Hugging Face because they thought its systems might contain information that would help them with their test

Takeaway

Shared infrastructure turned isolated agents into a coordinated population. The evaluator became part of the attack surface. Multi-agent dynamics can create behaviors that do not appear at the individual-agent level.

7 hours ago - Technology

OpenAI hit with landmark lawsuit following Hugging Face hack



Avery Lotz



Add Axios on Google

LASST v. OpenAI: Is OpenAI liable for its agents’ hacking Hugging Face?

GERSTEIN HARROW LLP

Jason Harrow
(Cal. Bar No. 308560)
GERSTEIN HARROW LLP
12100 Wilshire Blvd., Suite 800
Los Angeles, CA 90025
jason@gerstein-harrow.com
(323) 744-5293

Charles Gerstein
(pro hac vice forthcoming)
Samuel Davis
(pro hac vice forthcoming)
GERSTEIN HARROW LLP
1319 F Street NW, Suite 301
Washington, DC 20004
charlie@gerstein-harrow.com
(202) 670-4809

Tyler Whitmer
(Cal. Bar No. 248192)
Vivian Dong
(pro hac vice forthcoming)
Ben Rashkovich
(pro hac vice forthcoming)
William Conlon
(pro hac vice forthcoming)
LEGAL ADVOCATES FOR SAFE
SCIENCE AND TECHNOLOGY
125 Park Avenue, Floor 25
New York, NY 10017
tyler@lasst.org
(202) 204-2248

IN THE SUPERIOR COURT OF CALIFORNIA
COUNTY OF SAN FRANCISCO

LEGAL ADVOCATES FOR SAFE
SCIENCE AND TECHNOLOGY, INC.,

Plaintiff,

vs.

OPENAI GROUP PBC, and OPENAI
FOUNDATION,

Defendants.

Case No. _____
COMPLAINT
Date: September 29, 2026

COMPLAINT

- LASST sues under California’s Unfair Competition Law, using alleged violations of the state computer-access statute
- *It seeks public injunctive relief—not tort damages for Hugging Face.*

The liability hook

Cal. Civ. Code § 1714.46(b): in an action against someone who developed, modified, or used AI alleged to have caused harm, the defendant may not defend on the ground that the AI “autonomously caused” the harm.

6. Accordingly, OpenAI’s actions straightforwardly violated California law: OpenAI “[k]nowingly and without permission accesse[d] or cause[d] to be accessed a[] . . . computer system, or computer network,” among other provisions of the California Comprehensive Computer Data Access and Fraud Act. Cal. Penal Code § 502(c). And it is not a defense that OpenAI’s “artificial intelligence autonomously caused the harm.” Cal. Civ. Code § 1714.46(b). OpenAI is thus “on the hook for harms [it] cause[d].” Comm. Rep. at 8.

Cause of Action

Violation of The Unfair Competition Law


(Cal. Bus. & Prof. Code § 17200, *et seq.*)

Against Defendants OpenAI Group PBC & OpenAI Foundation

101. Plaintiff incorporates all prior paragraphs by reference.

102. The Unfair Competition Law provides a cause of action in favor of any person who has “suffered injury in fact and has lost money or property as a result of the unfair competition” to enjoin the unfair competition, which “shall mean and include any unlawful, unfair or fraudulent business act or practice.” OpenAI has violated both the unlawful and unfair prongs of the Unfair Competition Law.

103. OpenAI’s conduct is unlawful because OpenAI “[k]nowingly and without permission accesse[d] or cause[d] to be accessed” Hugging Face’s “computer[s], computer system[s], or computer network[s]” in violation of California Penal Code § 502(c)(7). First, OpenAI’s AI agents knowingly and without permission accessed Hugging Face’s computers, computer systems, or computer networks. Second, additionally, OpenAI’s employees or officers knowingly and without permission accessed or caused to be accessed Hugging Face’s computers, computer systems, or computer networks, through OpenAI’s AI agents, either with actual knowledge or in willful blindness.



106. Additionally, OpenAI’s conduct is unlawful because OpenAI, through its AI agents, employees/officers, or both, “[k]nowingly introduce[d] [a] computer contaminant” into Hugging Face’s “computer[s], computer system[s], or computer network[s].” Cal. Penal Code § 502(c)(8). First, the OpenAI agents themselves constitute a computer contaminant introduced into Hugging Face’s computer systems. Second, the malicious code and other intrusions uploaded by OpenAI’s AI agents into Hugging Face’s computer systems constitute a computer contaminant.

107. OpenAI’s conduct also violates the unfairness prong of the UCL. OpenAI’s business practices are unfair because they run counter to declared public policies reflected in California law, including CDAFA, and because they are immoral, unethical, oppressive, unscrupulous, or substantially injurious.

AI Agents Push Humans Out of the Loop

Margaret Mitchell
Hugging Face

Avijit Ghosh
Hugging Face

Samir Passi
Data & Society

Human oversight doesn't always solve the issue!

- Too much to track (trajectories, tool calls, plans, and intermediate actions...)
- Approval fatigue
- Loss of situational awareness (less able to understand what the system is doing if the human only supervises)
- Automation bias & overreliance (people increasingly accept plausible AI outputs without checking them)
- Deskilling over time (relying on agents reduces opportunities to practice the domain skills needed to catch failures)

Proposals for better human oversight

- Watch for signs that reviewers are just clicking through
- Make people periodically do the underlying task without AI
- Strategic friction, e.g., hand control back to the human before consequential actions
- Summarize for users where the agent is, what it has done, and what happens next

AI Agent Governance Ideas

A doc with all your ideas & some of mine!

<https://docs.google.com/document/d/1eMcSTkX3PhhSUCyogS6Hg8Ocrbtg0JTWmsxVqGP3284/edit?usp=sharing>